

Performance issues

Indexes

- Indexes are a common way to **enhance database performance**.
 - An index allows the database server to find and retrieve specific rows **much faster** than it could do without an index.
 - But indexes also **add overhead to the database system** as a whole, so they should be used sensibly

Create index

- `CREATE INDEX test1_id_index ON test1 (id);`
 - `CREATE INDEX test1_id_index ON test1 USING btree (id);`
 - `CREATE INDEX test1_id_index ON test1 [USING btree] (id) WHERE <condition>;`
- ➔ Partial index

Index types in PostgreSQL

- PostgreSQL provides several index types: B-tree, Hash, GiST, SP-GiST, GIN
- Each index type uses a different algorithm that is best suited to different types of queries.

By default, the CREATE INDEX command creates **B-tree** indexes, which **fit the most common situations**

Index types in PostgreSQL

- B-Tree (default)
 - handle equality and range queries on data that can be sorted into some ordering.
 - Operators: $<$, \leq , $=$, \geq , $>$, LIKE (col LIKE 'foo%' but not col LIKE '%bar')
 - Sorted output
- Hash index: can only handle simple equality comparisons
- GiST index: for several two-dimensional geometric data types,
 - not a single kind of index, but rather an infrastructure within which many different indexing strategies can be implemented
- GIN index
 - inverted indexes which can handle values that contain more than one key, arrays for example

Multicolumn index

- `CREATE INDEX test2_mm_idx ON test2 (major, minor);`
- B-Tree
- GiST index
- GIN index

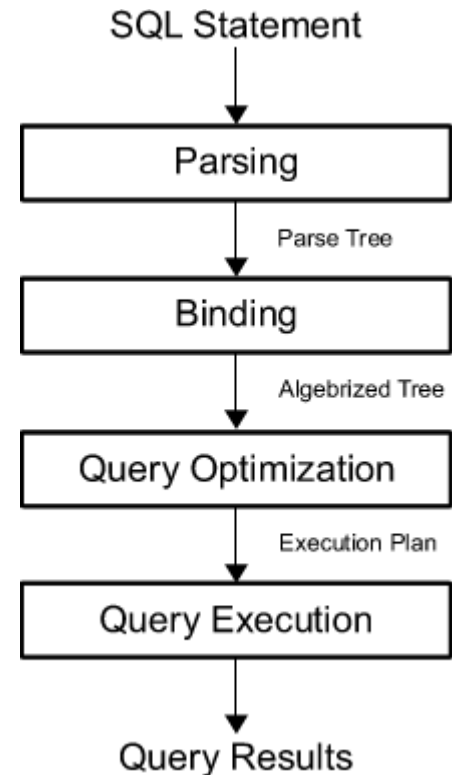
<https://www.postgresql.org/docs/10/sql-createindex.html>

<https://www.postgresql.org/docs/10/indexes.html>

Examining index usage

- **EXPLAIN [ANALYZE] [VERBOSE] *statement***
 - **EXPLAIN *statement***: displays the execution plan that the PostgreSQL planner generates for the supplied statement.

Actually two numbers are shown: the **start-up cost** before the first row can be returned, and the **total cost** to return all the rows.
 - **VERBOSE** option: displays additional information regarding the plan (output column list, table and function names, ...)



Examining index usage

- **EXPLAIN [ANALYZE] [VERBOSE] *statement***
 - **ANALYZE** option: causes the statement to be **actually executed**, not only planned, actual runtime statistics are added to the display
- **Important:** If you wish to use **EXPLAIN ANALYZE** on an INSERT, UPDATE, DELETE, CREATE TABLE AS, or EXECUTE statement without letting the command affect your data, use this approach:

BEGIN;

EXPLAIN ANALYZE ...;

ROLLBACK;

View table indexes

- \d table_name
- Ex.: \d customers

Tips

- Select **fewer columns** to improve hash join performance
- Index the *independent* **where** predicates to improve hash join performance

Tips

- Having a **WHERE / HAVING** clause in your queries does not necessarily means that it is a bad query
- Only retrieve the data you need
 - remove unnecessary columns from **SELECT**
 - **Inner join** vs. exists (with subqueries)
 - Select **DISTINCT** : try to avoid if you can
 - **LIKE** operator: the **index isn't used** if the pattern **starts with % or _**

- Limit your results : LIMIT, TOP
- Don't Make Queries More Complex Than They Need To Be
 - OR / IN / UNION ?
 - OR operator : index is not used except composite index → IN/UNION/OUTER JOIN
 - NOT operator: index is not used => avoid
 - AND vs BETWEEN
 - ANY / ALL: index not used => max , min ,...
 - Isolate columns in Condition : $\text{age} + 7 < 20 \rightarrow \text{age} < 13$

- Limit your results : **LIMIT, TOP**

You can add the LIMIT or TOP clauses to your queries to set a maximum number of rows for the result set.

```
SELECT TOP 3 *  
FROM customers;
```

```
SELECT *  
FROM customers  
LIMIT 3;
```

- Don't Make Queries More Complex Than They Need To Be
 - OR / IN / UNION ?
 - OR operator : index is not used except composite index → IN/UNION/OUTER JOIN
 - Using a condition with IN or UNION:

```
SELECT * FROM orderlines
WHERE orderid = 1 OR orderid = 5; -- (first cost: 8 - total cost: 4)
```

```
SELECT * FROM orderlines
WHERE orderid IN (1,5); -- (0.29 - 30)
```

```
SELECT * FROM orderlines
WHERE orderid = 1
UNION
SELECT * FROM orderlines
WHERE orderid = 5; -- (30 - 31)
```

- Don't Make Queries More Complex Than They Need To Be
 - To be **careful not to unnecessarily use** the UNION operation because you go through the same table multiple times → use a UNION in your query, the execution time will increase.
 - Alternatives to the UNION operation are: reformulating the query in such a way that all conditions are placed in one SELECT instruction, or **using an OUTER JOIN** instead of UNION.

```
SELECT P.* , o.quantity
FROM products p left join orderlines o ON(p.prod_id =
o.prod_id) -- (326 - 2076), ~500ms
WHERE o.orderlineid IS NULL; -- (326 - 2076), 162ms
```

```
SELECT * , 0
FROM products
WHERE prod_id not in (select prod_id from orderlines)
UNION
SELECT p.*, quantity
FROM products p join orderlines o ON(p.prod_id = o.prod_id);
-- (17 780 - 19 210) 864 ms
```


- **NOT operator**: index is not used => avoid

```
select * from customers  
where customerid != 5000;
```

```
select * from customers  
where customerid = 5000;
```

- **ANY / ALL**: index not used => max , min ,...

- Isolate columns in Condition :

age + 7 < 20 → **age < 13**

- No Brute force
 - JOIN clause:
 - Order of tables => biggest table: placed last in join
 - No redundant conditions on joins
 - Having clause:
 - Used only if needed
 - Not to replace WHERE => WHERE help to limit the intermediate number of records
- ➔ Need smart indexing, smart using

Lưu ý: có thể tìm thấy

- Geometric type :
 - <https://www.postgresql.org/docs/10/datatype-geometric.html>
 - <https://www.postgresql.org/docs/10/functions-geometry.html>
- GiST:
<https://www.postgresql.org/docs/10/indexes-types.html>