



Database

Lesson 6. Constraints and Triggers

Viet-Trung Tran

<https://is.hust.edu.vn/~trungtv/>

Learning Map

Sequence	Title
1	Introduction to Databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

Outline

- Overview
- Constraints
- Triggers

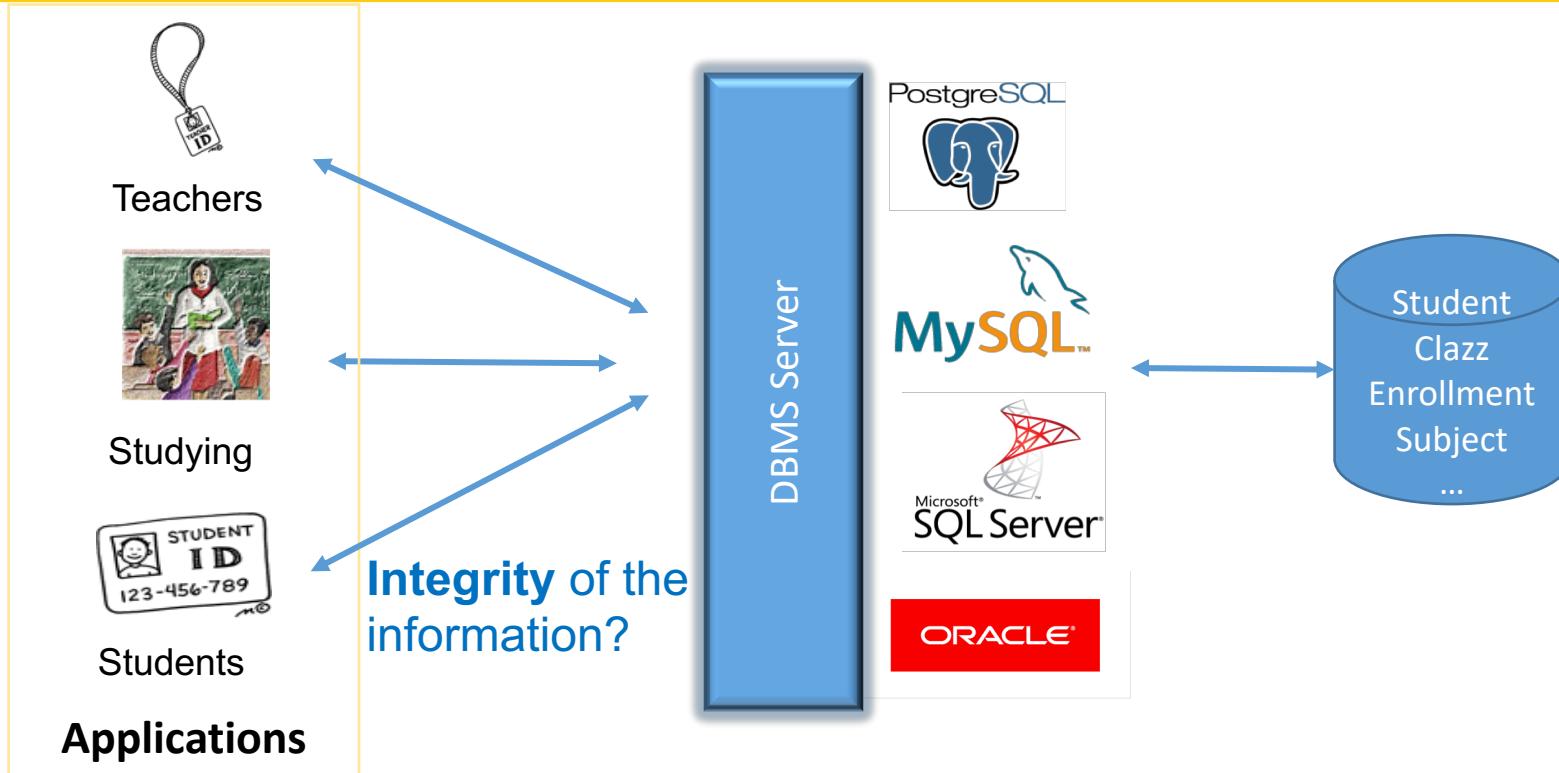
Learning objectives

- Upon completion of this lesson, students will be able to:
 - Well known about different constraints and define them correctly
 - Understand triggers: What is a trigger? how it works ? When using?
 - Define simple triggers

Keywords

Keyword	Description
Constraints	Constraints are the rules enforced on the data columns of a table. Constraints could be either on a column level or a table level
Triggers	A trigger is a SQL procedure that initiates an action (i.e., fires an action) when an event (INSERT, DELETE or UPDATE) occurs. They are stored in and managed by the DBMS
PL/SQL	Procedural Language/Structured Query Language is Oracle Corporation's procedural extension for SQL and the Oracle relational database

1. Overview



1. Overview: Database Schema

student(student_id, first_name, last_name, **dob**, **gender**,
address, note, **email**, **clazz_id**)

clazz(clazz_id, name, **lecturer_id**, **monitor_id**, **number_students**)

subject(subject_id, name, **credit**, percentage_final_exam)

enrollment(student_id, subject_id, **semester**, **midterm_score**, **final_score**)

lecturer(lecturer_id, first_name, last_name, dob, **gender**, address, email)

teaching(subject_id, lecturer_id)

grade(code, from_score, to_score)

1. Overview : Constraints and Triggers

- A **constraint** is a relationship among data elements that the DBMS is required to enforce
 - Example: key constraints
- **Triggers** are only executed when a specified condition occurs, e.g., insertion of a tuple
 - Easier to implement than complex constraints

2. Constraints: Kinds of Constraints

- Keys: PRIMARY KEY vs. UNIQUE
- Foreign-key, or referential-integrity
- Attribute-based constraints
 - Constrain values of a particular attribute.
- Tuple-based constraints
 - Relationship among components
- Assertions: any SQL Boolean expression

2.1. Keys: PRIMARY KEY vs. UNIQUE

- Declaring: similar syntax as primary key
- Example:

```
CREATE TABLE student (
    student_id CHAR(8) NOT NULL,
    first_name VARCHAR(20) NOT NULL,
    last_name VARCHAR(20) NOT NULL,
    ...
    email varchar(50) UNIQUE,
    clazz_id CHAR(8),
    CONSTRAINT student_pk PRIMARY KEY (student_id));
```

2.1. Keys: PRIMARY KEY vs. UNIQUE (2)

	PRIMARY KEY	UNIQUE KEY
Number defined on table	One	Multiple
Null columns allowed	No	Yes
Default index	CLUSTERED	NON-CLUSTERED
Purpose	Enforce Entity Integrity	Enforce Unique Data
Number of columns	One or more columns	One or more columns
Referenced by a Foreign Key Constraint	Yes	Yes

2.2. Foreign keys: Expressing Foreign Keys

- Use keyword REFERENCES, either:
 - After an attribute (for one-attribute keys)
 - As an element of the schema:
[CONSTRAINT <name>] FOREIGN KEY (<list of attributes>
REFERENCES <relation> (<attributes>))
- Referenced attributes must be declared PRIMARY KEY or UNIQUE

2.2. Foreign keys: Example

```
CREATE TABLE clazz (
    clazz_id CHAR(8) NOT NULL PRIMARY KEY,
    name VARCHAR(20), ... );

CREATE TABLE student (
    student_id CHAR(8) NOT NULL,
    ...
    clazz_id CHAR(8),
    CONSTRAINT student_pk PRIMARY KEY (student_id));

ALTER TABLE student ADD CONSTRAINT student_fk_class
FOREIGN KEY (clazz_id) REFERENCES clazz(clazz_id);
```

2.2. Foreign keys: Enforcing constraint

- An insert or update to student that introduces a non-existent `clazz_id` (`clazz_id` value is not found in `clazz`)
→ Reject
- A deletion or update to `clazz` that removes a `clazz_id` value found in some tuples of `student`?
 - Default: reject the modification
 - Cascade: make the same changes in `student`
 - Set NULL: change `clazz_id` in `student` to NULL

2.2. Foreign keys: Choosing policy

```
ALTER TABLE student
    ADD CONSTRAINT student_fk_class FOREIGN KEY
(clazz_id) REFERENCES clazz(clazz_id)
    ON DELETE SET NULL
    ON UPDATE CASCADE;
```

2.3. Attribute-based checks: Declaring

- Constraints on the value of a particular attribute
 - Add **CHECK(<condition>)** to the declaration for the attribute or add as relation-schema element
 - The condition may use the name of the attribute, but **any other relation or attribute name must be in a subquery**
- Example:

```
CREATE TABLE student (
    student_id CHAR(8) NOT NULL PRIMARY KEY, ...,
    gender CHAR(1),
    clazz_id CHAR(8) CHECK (clazz_id IN (SELECT clazz_id FROM clazz)),
    CONSTRAINT student_chk_gender CHECK (gender = 'F' OR gender = 'M') );
```

2.3. Attribute-based checks: Timing of checks

- Only when a value for that attribute is inserted or updated

```
CREATE TABLE student (
    student_id CHAR(8) NOT NULL PRIMARY KEY, ...,
    gender CHAR(1),
    clazz_id CHAR(8) CHECK (clazz_id IN (SELECT
        clazz_id FROM clazz)),
    CONSTRAINT student_chk_gender CHECK (gender = 'F'
    OR gender = 'M')) ;
```

Not checked if a class is deleted
from **clazz**

2.4. Tuple-based checks

- **CHECK (<condition>)** may be added as a relation-schema element.
- The condition may **refer to any attribute** of the relation
 - But other attributes or relations require a subquery
- Timing of checks: **on insert or update only.**

```
CREATE TABLE grade(
    code CHAR(1) NOT NULL,
    from_score DECIMAL(3,1) NOT NULL,
    to_score DECIMAL(3,1) NOT NULL, ...,
    CONSTRAINT grade_chk_toScore CHECK (to_score >
from_score) );
```

2.5. Assertions: Declaring

- Database-schema elements, like relations or views
- Defined by:
`CREATE ASSERTION <name>
CHECK (<condition>);`
- Condition may refer to any relation or attribute in the database schema
- Drop an assertion:
`DROP ASSERTION <assertion name>;`

2.5. Assertions: Example

```
CREATE ASSERTION teachingSubject CHECK (
    (SELECT COUNT(*) FROM teaching) >=
    (SELECT COUNT(*) FROM subject) );
```

```
CREATE ASSERTION numberStdInClass CHECK (
    NOT EXISTS (
        SELECT * FROM clazz c
        WHERE number_students <>
            (SELECT count(*) FROM student
            WHERE clazz_id = c.clazz_id)
    )
);
```

2.5. Assertions: Timing of Assertion Checks

- In principle, we must check every assertion **after every modification** to any relation of the database
- A **clever system** can observe that only certain changes could cause a given assertion to be violated
 - No change to student can affect teaching Subject
 - Neither can an insertion to teaching
- Very **hard to implement** assertions efficiently

3. Triggers

- Motivation
- Trigger Syntax
- Using triggers
- Examples

3.1 Trigger motivation

- Assertions
 - powerful,
 - but the DBMS often can't tell when they need to be checked
- Attribute and tuple-based checks
 - checked at known times,
 - but are not powerful in some circumstances
- Triggers let the user decide when to check for any condition

3.1. Motivation: ECA Rules

- A trigger defines an operation that is performed when a **specific event occurs on a relation**:
 - inserts a new record / updates an existing record / deletes a record
- Trigger functions have access to special variables from the database engine
- Called also **ECA rules (Event-Condition-Action)**
 - Event: type of database modification
 - Condition: Any SQL Boolean-valued expression
 - Action: Any SQL statements

3.1. Motivation: Example

- Constraint: when a new student is inserted into student relation, the number of students in his class must be increased
 - student(student_id, first_name, last_name, dob, gender, address, note, email, **clazz_id**)
 - clazz(clazz_id, name, lecturer_id, monitor_id, **number_students**)

The diagram illustrates the structure of a MySQL trigger definition. It consists of three main parts: Event, Condition, and Action.

Event: A green bracket covers the first two lines of the code: `CREATE TRIGGER clazz_changes_tg` and `AFTER INSERT ON student`. A green arrow points from this bracket to the word "Event".

Condition: A purple bracket covers the line `WHEN (nnn.clazz_id IS NOT NULL)`. A purple arrow points from this bracket to the word "Condition".

Action: A blue bracket covers the entire `BEGIN` block, which contains the `update` statement. A blue arrow points from this bracket to the word "Action".

```
CREATE TRIGGER clazz_changes_tg
AFTER INSERT ON student
REFERENCING NEW ROW AS nnn
FOR EACH ROW
WHEN (nnn.clazz_id IS NOT NULL)
BEGIN
    update clazz
    set number_students = number_students + 1
    where clazz_id = nnn.clazz_id;
END;
```

3.2. Trigger syntax

- Creating a trigger:

```
CREATE [OR REPLACE] TRIGGER <trigger_name>
    {BEFORE | AFTER | INSTEAD OF }
    {INSERT | DELETE | UPDATE [OF <attribute_name>] }
    ON <table_name>
    REFERENCING {NEW | OLD} {ROW | TABLE} AS <name>
    [FOR EACH ROW ]
    [WHEN (<condition>) ]
    BEGIN
        <trigger body goes here >
    END;
```

- Dropping a trigger:

```
DROP TRIGGER <trigger_name>;
```

3.2. Trigger syntax: Event

- AFTER, BEFORE, INSTEAD OF:
 - AFTER, BEFORE: used for tables / views
 - INSTEAD OF: used only for views
 - A way to execute view modifications: triggers translate them to appropriate modifications on the base tables
- INSERT, DELETE, UPDATE , UPDATE OF
 - UPDATE OF <columns>: update on a particular column

3.2. Trigger syntax: Level

- Row-level trigger:
 - Indicated by option FOR EACH ROW
 - Trigger executes once for each modified tuple
- Statement-level trigger:
 - Without option FOR EACH ROW or with FOR EACH STATEMENT
 - Trigger execute once for a SQL statement, regardless how many tuples are modified

3.2. Trigger syntax: REFERENCING

- INSERT statements imply a new tuple (for row-level) or new table (for statement-level)
 - The table is the set of inserted tuples
- DELETE implies an old tuple or table
- UPDATE implies both
- Refer to these by
 - **REFERENCING [NEW | OLD] [TUPLE | TABLE] AS <name>**
- **Each DBMS has its own implementation**, REFERENCING may not be used in:
 - Access directly to special variables from the database engine: NEW, OLD,...

3.2. Trigger syntax: Condition

- Any boolean-valued condition
- Evaluated on the database as it would exist **before or after the triggering event**, depending on whether BEFORE or AFTER is used.
 - But always before the changes take effect.
- Access the new/old tuple/table through the names in the REFERENCING clause

Trigger syntax: Action

- Can be more than one SQL statement:
 - Surrounded by BEGIN .. END
- Language:
 - Simple SQL statements
 - Extension of SQL: procedural languages, depends on each DBMS
 - PL/SQL (Oracle), PL/pgSQL (PostgreSQL), T-SQL(SQL Server) ,..

3.3. Using triggers: When

- Auditing data modification (keeping history of data), providing transparent event logging
- Validation and business security checking if so is desired
 - Eg. column formatting before and after inserts into database
- Enforcing complex integrity constraints
- Enforcing complex business rules
- Maintaining replicate tables
- Building complex views that are updatable

3.3. Using triggers: Guidelines for designing triggers

- Do not define triggers that duplicate database features
 - do not define triggers to reject bad data if you can do the same checking through constraints
- Use triggers **only for centralized, global operations** that must fire for the triggering statement, regardless of which user or database application issues the statement
- Do not create recursive triggers
- Use triggers on DATABASE judiciously (e.g. server error, logon, logoff,...):
 - they are executed for every user every time the event occurs on which the trigger is created

3.4. Trigger example: In Oracle

- Add a new column in clazz relation

```
alter table clazz  
add column number_students integer not null default 0;
```

- Create a trigger on student relation

```
CREATE TRIGGER clazz_changes_tg  
AFTER UPDATE ON student  
FOR EACH ROW  
WHEN (:NEW.clazz_id <> :OLD.clazz_id)  
BEGIN  
    update clazz set number_students= number_students+1  
    where clazz_id = :NEW.clazz_id;  
    update clazz set number_students = number_students-1  
    where clazz_id = :OLD.clazz_id;  
END;
```

3.4. Trigger example: In PostgreSQL

```
CREATE FUNCTION public.tg_fnc_change_clazz()
    RETURNS trigger LANGUAGE 'plpgsql' AS $$

BEGIN
    update clazz set number_students = number_students+1
        where clazz_id = NEW.clazz_id;
    update clazz set number_students = number_students-1
        where clazz_id = OLD.clazz_id;
    return NEW;
END; $$

CREATE TRIGGER tg_af_update_clazz
AFTER UPDATE OF clazz_id
ON student
FOR EACH ROW
EXECUTE PROCEDURE public.tg_fnc_change_clazz();
```

Remark

- Constraints, Assertions, Triggers:
 - How to declare
 - Timing of checks
 - Differences
 - Only use them if you really need to, especially triggers
 - Each DBMS has its own variation in implementation:
 - Options
 - Syntax: triggers as an example
- ➔ Reading documentation for each DBMS installed

Quiz 1.

Quiz Number	1	Quiz Type	OX	Example Select
Question	How many attributes are there in one unique key?			
Example	<ul style="list-style-type: none">A. Only oneB. One, two or more			
Answer	B			
Feedback	As other constraints, each constraint can have many attributes			

Quiz 2.

Quiz Number	2	Quiz Type	OX	Example Select
Question	A trigger can be explicitly called by user?			
Example	<ul style="list-style-type: none">A. YesB. No			
Answer	B			
Feedback	No, trigger is invoked by an event occurring inside database engine. You can not explicitly invoke it			

Quiz 2.

Quiz Number	3	Quiz Type	OX	Example Select
Question	Application users know while a trigger is executed?			
Example	<ul style="list-style-type: none">A. YesB. No			
Answer	B			
Feedback	Trigger is invoked by database engine, it is transparent to users			

Summary

- Introduction
 - Why we need constraints and triggers?
- Constraints
 - Keys
 - Foreigner key
 - Check constraints
 - Assertion
- Triggers
 - Motivation
 - Triggers definition
 - Using triggers



Next lesson: Entity Relationship Model

- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom. Database Systems: The Complete Book. Pearson Prentice Hall. the 2nd edition. 2008: Chapter 7
- Nguyen Kim Anh, Nguyên lý các hệ cơ sở dữ liệu, NXB Giáo dục. 2004: Chương 7