



Database

Lesson 11. Query processing

Viet-Trung Tran

<https://is.hust.edu.vn/~trungtv/>

Learning Map

Sequence	Title
1	Introduction to Databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

Outline

- Overview
 - What is query processing
 - Phrases of query processing
 - Parser
 - Optimizer
- Understanding query optimizer
 - Step 1: Equivalence transformation
 - Step 2: Annotation for the algorithm of the RA expression
 - Step 3: Cost estimation for different query execution plans

Objectives

- Upon completion of this lesson, students will be able to:
 - Understanding different phases of query processing
 - Understanding the implementation of query optimizer

Keywords

Query processing	Activities involved in retrieving/storing data from/to the database
Query optimization	Selection of an efficient query execution plan
Relational algebra	An algebra whose operands are relations or variables that represent Relations

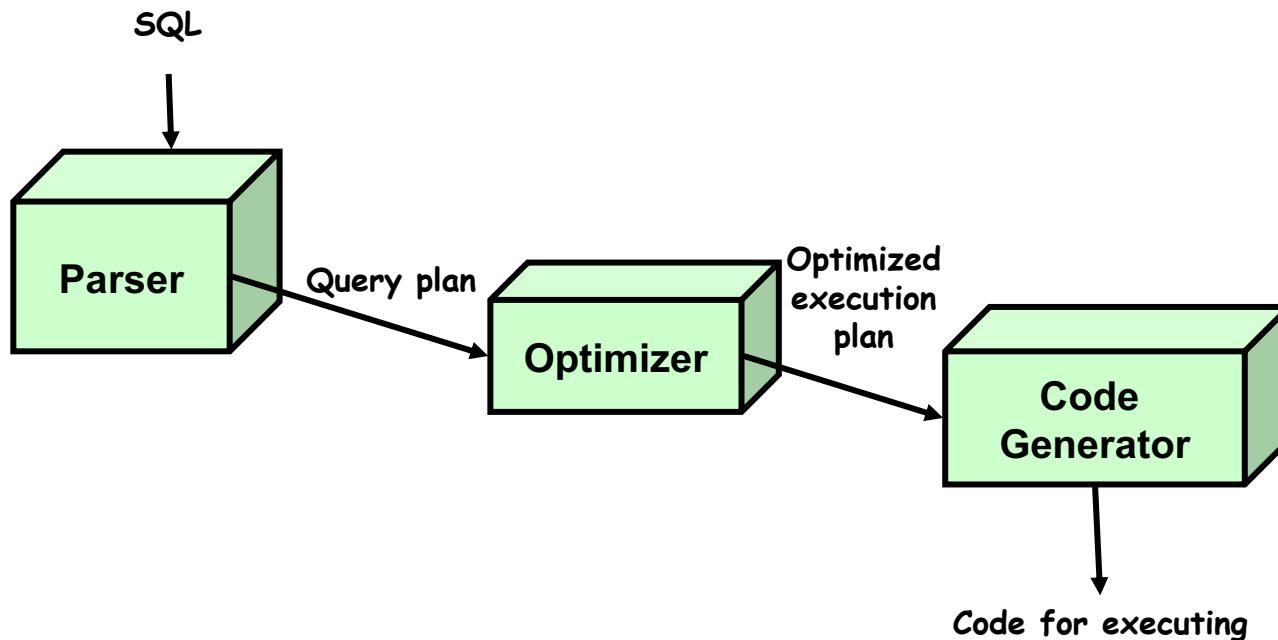
1. Overview

- What is query processing
- Phrases of query processing
- Parser
- Optimizer

1.1. What is query processing

- The entire process or activities involved in retrieving data from the database
 - SQL query translation into low level instructions (usually relational algebra)
 - Query optimization to save resources, cost estimation or evaluation of query
 - Query execution for the extraction of data from the database.

1.2. Phases of query processing



1.3. Parser

- Scans and parses the query into individual tokens and examines for the correctness of query
 - Does it contain the right keywords?
 - Does it conform to the syntax?
 - Does it contain the valid tables, attributes?
- Output: Query plan
 - E.g.
 - Input: $\text{SELECT balance FROM account WHERE balance < 2500}$
 - Output: Relational algebra expression $\sigma_{balance < 2500}(\Pi_{balance}(account))$
 - But it's not unique $\Pi_{balance}(\sigma_{balance < 2500}(account))$

1.4. Optimizer

- Input: RA expression
- Output: Query execution plan
- Query execution plan = query plan + the algorithms for the executions of RA operations
- Aims to choose the cheapest execution plan out of the possible ones
 - Step 1: Equivalence transformation
 - Step 2: Annotation for the algorithm of the RA expression
 - Step 3: Cost estimation for different query execution plans

$$\Pi_{balance}(\sigma_{balance < 2500}(account))$$

$\Pi_{balance}$
|
 $\sigma_{balance < 2500}$
use index 1
|
account

2. Understanding optimizer

- Choose the cheapest execution plan out of the possible ones
 - Step 1: Equivalence transformation
 - Step 2: Annotation for the algorithmic execution of the RA expression
 - Step 3: Cost estimation for different query execution plans

2.1. Step 1: Equivalence transformation

- RA expressions are equivalent if they generate the same set of tuples on every database instance
- **Equivalence rules:**
 - Transform one relational algebra expression into equivalent one
 - Similar to numeric algebra: $a + b = b + a$, $a(b + c) = ab + ac$, etc
- **Why producing equivalent expressions?**
 - equivalent algebraic expressions give the same result
 - but usually the execution time varies significantly

2.1. Step 1: Equivalence transformation

- Equivalence transformation rules
- (1) Conjunctive selection operations can be deconstructed into a sequence of individual selections; cascade of σ
 - $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$
- (2) Selection operations are commutative
 - $\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$
- (3) Only the final operations in a sequence of projection operations is needed; cascade of Π
 - $\Pi_{L_1}(\Pi_{L_2}(\dots \Pi_{L_n}(E) \dots)) = \Pi_{L_1}(E)$
- (4) Selections can be combined with Cartesian products and theta joins
 - $\sigma_{\theta_1}(E_1 \times E_2) = E_1 \bowtie_{\theta_1} E_2$
 - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

2.1. Step 1: Equivalence transformation

- Equivalence transformation rules
- (5) Theta Join operations are commutative
 - $E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$
- (6) Natural join operations are associative
 - $E_1 \bowtie (E_2 \bowtie E_3) = (E_1 \bowtie E_2) \bowtie E_3$
 - Theta join are associative in the following manner where θ_2 involves attributes from E2 and E3 only
 - $(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$

2.1. Step 1: Equivalence transformation

- Equivalence transformation rules
- (7) Selection distributes over joins in the following ways
 - If predicate involves attributes of E1 only
 - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = \sigma_{\theta_1}(E_1) \bowtie_{\theta_2} E_2$
 - If predicate θ_1 involves only attributes of E1 and θ_2 involves only attributes of E2 (a consequence of rule 7 and 1)
 - $\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta_3} E_2) = \sigma_{\theta_1}(E_1) \bowtie_{\theta_3} \sigma_{\theta_2}(E_2)$

2.1. Step 1: Equivalence transformation

- Equivalence transformation rules
- (8) Projection distributes over join as follows
 - $\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = \Pi_{L_1}(E_1) \bowtie_{\theta} \Pi_{L_2}(E_2)$
 - If θ involves attributes in $L_1 \cup L_2$ only and L_i contains attributes of E_i
- (9) The set operations union and intersection are commutative
 - $E_1 \cup E_2 = E_2 \cup E_1$
 - $E_1 \cap E_2 = E_2 \cap E_1$
- (10) The union and intersection are associative
 - $(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$

2.1. Step 1: Equivalence transformation

- Equivalence transformation rules
- (11) The selection operation distributes over union, intersection, and set-difference
 - $\sigma_\theta(E_1 \cup E_2) = \sigma_\theta(E_1) \cup \sigma_\theta(E_2)$
 - $\sigma_\theta(E_1 \cap E_2) = \sigma_\theta(E_1) \cap \sigma_\theta(E_2)$
 - $\sigma_\theta(E_1 - E_2) = \sigma_\theta(E_1) - \sigma_\theta(E_2)$
- (12) The project operation distributes over the union
 - $\Pi_L(E_1 \cup E_2) = \Pi_L(E_1) \cup \Pi_L(E_2)$

2.2. Step 2: Execution algorithms of RA operations

- Algebra expression is not a query execution plan.
- Additional decisions required:
 - which indexes to use, for example, for joins and selects?
 - which algorithms to use, for example, sort-merge vs. hash join?
 - materialize intermediate results or pipeline them?

2.2. Step 2: Execution algorithms of RA operations

- Basic Operators
- One-pass operators:
 - Scan
 - Select
 - Project
- Multi-pass operators:
 - Join
 - Various implementations
 - Handling of larger-than-memory sources
 - Aggregation, union, etc.

2.2. Step 2: Execution algorithms of RA operations

- 1-Pass Operators: Scanning a Table
 - Sequential scan: read through blocks of table
 - Index scan: retrieve tuples in index order

2.2. Step 2: Execution algorithms of RA operations

- Nested-loop JOIN

```
For each tuple tr in r {  
    for each tuple ts in s {  
        if (tr and ts satisfy the join condition) {  
            add tuple tr x ts to the result set  
        }  
    }  
}
```

- No index needed
- Any join condition types
- Expensive: $O(n^2)$

2.2. Step 2: Execution algorithms of RA operations

- Single-loop JOIN (Index-based)

```
for each tube tr in R {  
    seach for ts in s thought index {  
        if ts.exist() {  
            add tr x ts to the result set  
        }  
    }  
}
```

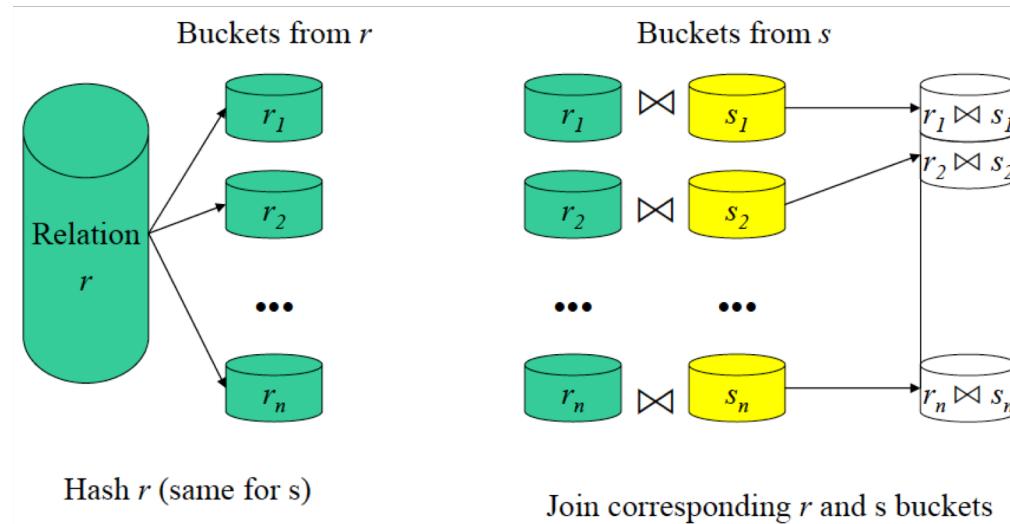
- Index needed
- Cheaper: $O(n \log m)$

2.2. Step 2: Execution algorithms of RA operations

- Sort-merge JOIN
 - Requires data physically sorted by join attributes: Merge and join sorted files, reading sequentially a block at a time
 - Maintain two file pointers
 - While tuple at R < tuple at S, advance R (and vice versa)
 - While tuples match, output all possible pairings
 - Very efficient for presorted data. Otherwise, may require a sort (adds cost + delay)

2.2. Step 2: Execution algorithms of RA operations

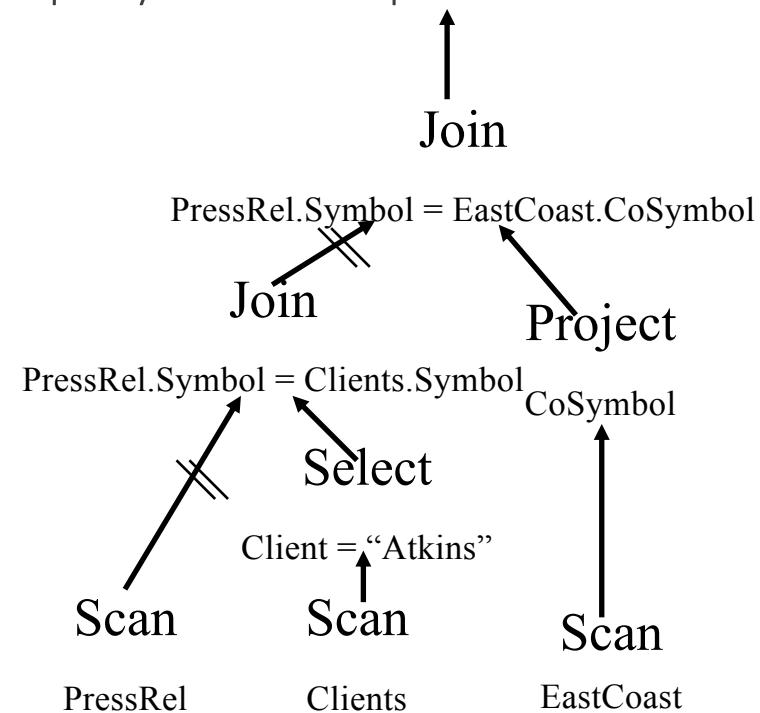
- Partition-hash JOIN
 - Hash two relations on join attributes
 - Join buckets accordingly



2.2. Step 2: Execution algorithms of RA operations

- Execution Strategy: Materialization vs. Pipelining

- Execution strategy defines how to walk the query execution plan
 - Materialization
 - Pipelining



2.2. Step 2: Execution algorithms of RA operations

- Materialization
 - Performs the innermost or leaf-level operations first of the query execution plan
 - The intermediate result of each operation is materialized into temporary relation and becomes input for subsequent operations.
 - The cost of materialization is the sum of the individual operations plus the cost of writing the intermediate results to disk
 - lots of temporary files, lots of I/O.

2.2. Step 2: Execution algorithms of RA operations

- **Pipelining**

- Operations form a queue, and results are passed from one operation to another as they are calculated
- Pipelining restructures the individual operation algorithms so that they take streams of tuples as both input and output.
- **Limitation**

- algorithms that require sorting can only use pipelining if the input is already sorted beforehand
- since sorting by nature cannot be performed until all tuples to be sorted are known.

2.3. Step 3: Cost estimation

- Each relational algebra expression can result in many query execution plans
- Some query execution plans may be better than others
- Finding the fastest one
 - Just an estimation under certain assumptions
 - Huge number of query plans may exist

2.3. Step 3: Cost estimation

- Cost estimation factors
 - Catalog information: database maintains statistics about relations
 - Ex.
 - number of tuples per relation
 - number of blocks on disk per relation
 - number of distinct values per attribute
 - histogram of values per attribute
 - Problems
 - cost can only be estimated
 - updating statistics is expensive, thus they are often out of date

2.3. Step 3: Cost estimation

- Choosing the cheapest query plan
 - Problem:
 - Estimating cost for all possible plans too expensive.
 - Solutions:
 - pruning: stop early to evaluate a plan
 - heuristics: do not evaluate all plans
 - Real databases use a combination of
 - Apply heuristics to choose promising query plans.
 - Choose cheapest plan among the promising plans using pruning.
 - Examples of heuristics:
 - perform selections as early as possible
 - perform projections early avoid Cartesian products

2.3. Step 3: Cost estimation

- Heuristic rules

- Break apart conjunctive selections into a sequence of simple selections
- Move σ down the query tree as soon as possible
- Replace σ - x pairs by \bowtie
- Break apart and move Π down the tree as soon as possible
- Perform the joins with the smallest expected result first

Remark

- Query processing is the entire process or activities involved in retrieving data from the database
 - Parser
 - Optimizer
 - Code generator
- Query optimizer
 - Step 1: Equivalence transformation
 - Step 2: Annotation for the algorithm of the RA expression
 - Step 3: Cost estimation for different query execution plans

Quiz 1.

Quiz Number	1	Quiz Type	OX	Example Select
Question	What is the output of the parser in query processing?			
Example	<ul style="list-style-type: none">A. Query execution planB. Query planC. Relational algebra expressionD. Code for executing			
Answer	B and C			
Feedback	The output of the parser is a query plan, usually in term of a relational algebra expression			

Quiz 2.

Quiz Number	2	Quiz Type	OX	Example Select
Question	What can the query optimizer do?			
Example	<ul style="list-style-type: none">A. Equivalence transformationB. Annotation for the algorithmic execution of the RA expressionC. Cost estimationD. Executing the query plan and return results			
Answer	A, B and C			
Feedback	A, B and C are 3 common steps in query optimization			

Summary

- Overview of query processing
 - Parser
 - Optimizer
- Understanding query optimizer
 - Step 1: Equivalence transformation
 - Step 2: Annotation for the algorithmic execution of the RA expression
 - Step 3: Cost estimation for different query execution plans



Next lesson: Transaction management

- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom. Database Systems: The Complete Book. Pearson Prentice Hall. the 2nd edition. 2008: Chapter 7
- Nguyễn Kim Anh, Nguyên lý các hệ cơ sở dữ liệu, NXB Giáo dục. 2004: Chương 7