



TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Database

Lesson 12. Transaction management

Viet-Trung Tran

<https://is.hust.edu.vn/~trungtv/>

Learning Map

Sequence	Title
1	Introduction to Databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

Outline

- Transaction
 - Definition
 - ACID properties
 - Transaction states
 - Transaction management interfaces
- Concurrency control
 - Objective
 - Scheduling
 - Serializability
 - Locks
 - Locking protocol
 - 2 Phase Locking (2PL)

Objectives

- Upon completion of this lesson, students will be able to:
 - Understanding transaction and ACID properties
 - Understanding locking and 2 phase locking algorithm

1. Transaction

- Definition
- ACID properties
- Transaction states
- Transaction management interfaces

1.1. What is transaction

- A transaction is a unit of program execution that accesses and possibly updates various data items.
- Example: transfer \$50 from account A to account B
 - 1. $R(A)$
 - 2. $A \leftarrow A - 50$
 - 3. $W(A)$
 - 4. $R(B)$
 - 5. $B \leftarrow B + 50$
 - 6. $W(B)$
- Two main issues:
 - 1. concurrent execution of multiple transactions
 - 2. failures of various kind (e.g., hardware failure, system crash)

1.2. ACID Properties

- Database system must guarantee ACID for transactions:
 - **Atomicity**: either all operations of the transaction are executed or none
 - **Consistency**: execution of a transaction in isolation preserves the consistency of the database
 - **Isolation**: although multiple transactions may execute concurrently, each transaction must be unaware of the other concurrent transactions.
 - **Durability**: After a transaction completes successfully, changes to the database persist even in case of system failure.

1.2.1. Atomicity

- Example: transfer \$50 from account A to account B
 1. $R(A)$
 2. $A \leftarrow A - 50$
 3. $W(A)$
 4. $R(B)$
 5. $B \leftarrow B + 50$
 6. $W(B)$
- What if failure (hardware or software) after step 3?
 - money is lost
 - database is inconsistent
- **Atomicity:**
 - either all operations or none
 - updates of partially executed transactions not reflected in database

1.2.2. Consistency

- Example: transfer \$50 from account A to account B
 1. $R(A)$
 2. $A \leftarrow A - 50$
 3. $W(A)$
 4. $R(B)$
 5. $B \leftarrow B + 50$
 6. $W(B)$
- **Consistency in example:** sum $A + B$ must be unchanged
- **Consistency in general:**
 - explicit integrity constraints (e.g., foreign key)
 - implicit integrity constraints (e.g., sum of all account balances of a bank branch must be equal to branch balance)
- **Transaction:**
 - must see consistent database
 - during transaction inconsistent state allowed
 - after completion database must be consistent again

1.2.3. Isolation

- Example: transfer \$50 from account A to account B
 - 1. $R(A)$
 - 2. $A \leftarrow A - 50$
 - 3. $W(A)$
 - 4. $R(B)$
 - 5. $B \leftarrow B + 50$
 - 6. $W(B)$
- **Imagine second transaction $T2$:**
 - $T2 : R(A), R(B), \text{print}(A + B)$
 - $T2$ is executed between steps 3 and 4
 - $T2$ sees an inconsistent database and gives wrong result


1.2.3. Isolation

- **Trivial isolation:** run transactions serially
- Isolation for concurrent transactions:
 - For every pair of transactions T_i and T_j ,
 - it appears to T_i as if either T_j finished execution before T_i started
 - or T_j started execution after T_i finished.

1.2.4. Durability

- When a transaction is done it commits.
- **Example:** transaction commits too early
 - transaction writes A, then commits
 - A is written to the disk buffer
 - then system crashes
 - value of A is lost
- **Durability:** After a transaction has committed, the changes to the database persist even in case of system failure.
- **Commit only after all changes are permanent:**
 - either written to log file or directly to database
 - database must recover in case of a crash

1.3. Transaction states



1.4. Transaction management interfaces

- Begin Trans
- Commit ()
- Abort()
- Savepoint Save()
- Rollback (savepoint)
(savepoint = 0 ==> Abort)

2. Concurrency control

- Objective
- Scheduling
- Serializability
- Locks
- Locking protocol
- 2 Phase Locking (2PL)

2.1. Objective

- Objective:
 - ensures that database transactions are performed concurrently without violating the data integrity
 - guarantees that no effect of committed transactions is lost, and no effect of aborted (rolled back) transactions remains in the related database.
- Example

<pre>T0: read(A); A := A - 50; write(A); read(B); B := B + 50; write(B);</pre>	<pre>T1: read(A); temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; write(B);</pre>
--	---

2.1. Scheduling

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(1)

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(2)

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(3)

2.2. Serializability

- A schedule of a set of transactions is a linear ordering of their actions
 - e.g. for the simultaneous deposits example:
R1(X) R2(X) W1(X) W2(X)
- A serial schedule is one in which all the steps of each transaction occur consecutively
- A serializable schedule is one which is equivalent to some serial schedule

2.3. Locks

- A lock is a mechanism to control concurrency on a data item.
 - one way of enforcing concurrency control policies
- Two types of locks on a data item A :
 - **exclusive** – $xL(A)$: data item A can be both read and written
 - **shared** – $sL(A)$: data item A can only be read.
- Lock request are made to **concurrency control manager**.
- Transaction is blocked until lock is granted.
- **Unlock** A – $uL(A)$: release the lock on a data item A

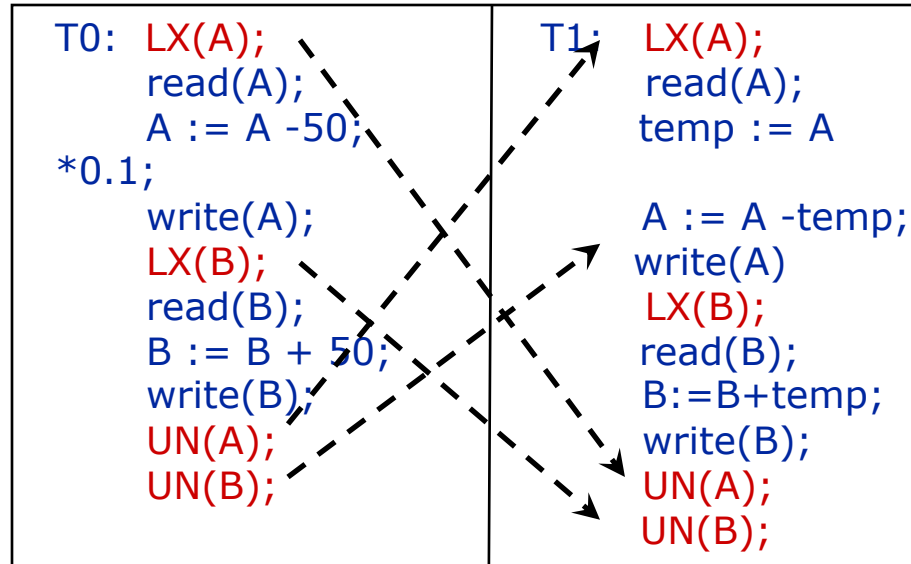
2.3.1. Lock Compatibility

- Lock compatibility matrix:

$T_1 \downarrow T_2 \rightarrow$	shared	exclusive
shared	true	false
exclusive	false	false

- T_1 holds shared lock on A:**
 - shared lock is granted to T_2
 - exclusive lock is not granted to T_2
- T_2 holds exclusive lock on A:**
 - shared lock is not granted to T_2
 - exclusive lock is not granted to T_2
- Shared locks can be shared by any number of transactions.**

2.3.2. Example

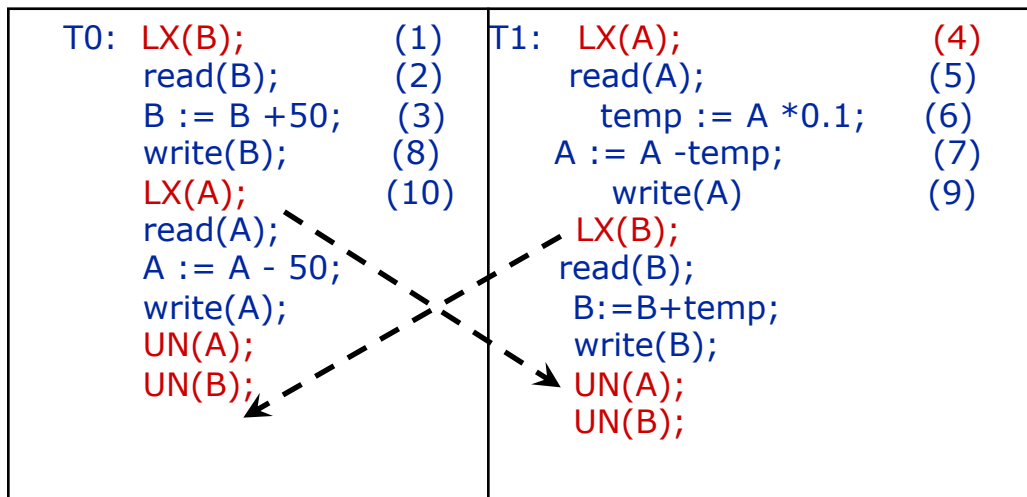


2.4. Locking protocol

- Example transaction T_2 with locking:
 1. $sL(A), R(A), uL(A)$
 2. $sL(B), R(B), uL(B)$
 3. $print(A + B)$
- **T_2 uses locking, but is not serializable**
 - A and/or B could be updated between steps 1 and 2
 - printed sum may be wrong
- **Locking protocol:**
 - set of rules followed by all transactions while requesting/releasing locks
 - locking protocol restricts the set of possible schedules

2.4.1. Pitfalls of locking protocols – Deadlock

- Example: two concurrent money transfers




- Deadlock: situation when transactions block each other
- Handling deadlocks:
 - one of the transactions must be rolled back (i.e., undone)
 - rolled back transaction releases locks

2.4.2. Pitfalls of locking protocols – Starvation

- **Starvation:** transaction continues to wait for lock
- **Examples:**
 - the same transaction is repeatedly rolled back due to deadlocks
 - a transaction continues to wait for an exclusive lock on an item while a sequence of other transactions are granted shared locks
- Well-designed concurrency manager avoids starvation.

2.5. 2 Phase Locking (2PL)

- Protocol that guarantees serializability.
 - Phase 1
 - locks are acquired and no locks are released
 - Phase 2
 - locks are released and no locks are acquired



2.5.1. Example

T_1
Lock(A)
Read(A)
Lock(B)
Read(B)
$B := B + A$
Write(B)
Unlock(A)
Unlock(B)

T_2
Lock(B)
Read(B)
Lock(A)
Read(A)
Unlock(B)
$A := A + B$
Write(A)
Unlock(A)

2PL

T_3
Lock(B)
Read(B)
$B = B - 50$
Write(B)
Unlock(B)
Lock(A)
Read(A)
$A = A + 50$
Write(A)
Unlock(A)

T_4
Lock(A)
Read(A)
Unlock(A)
Lock(B)
Read(B)
Unlock(B)
Print(A+B)

Not 2PL

Summary

- Transaction and ACID properties
- Concurrency control based on locking mechanisms
- 2 phase locking (2PL)

Next lesson: Transaction management

- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom. Database Systems: The Complete Book. Pearson Prentice Hall. the 2nd edition. 2008: Chapter 7
- Nguyen Kim Anh, Nguyên lý các hệ cơ sở dữ liệu, NXB Giáo dục. 2004: Chương 7